

Homework11

cat1997

December 8, 2019

1 Source files

1.1 centroids.c

```
#include <math.h>
#include <stdio.h>
#include <stdlib.h>

#include "data.h"
#include "clusters.h"

// perform kaverage clustering of data from file
double mpiStartTimer(){
    MPI_Barrier(MPI_COMM_WORLD);
    return MPI_Wtime();
}

double mpiEndTimer(double start){
    double end = MPI_Wtime();
    int N = 1;
    double *message = (double*) calloc(N, sizeof(double));
    double *elapsed = (double*) calloc(N, sizeof(double));
    message[0] = end-start;
    MPI_Allreduce(message, elapsed, 1, MPI_DOUBLE, MPI_MAX, MPI_COMM_WORLD);

    double res = elapsed[0];
    free(elapsed);
    free(message);

    return res;
}

int main(int argc, char **argv){
```

```

// initialize MPI
MPI_Init(&argc, &argv);

int size;
MPI_Comm_size(MPI_COMM_WORLD, &size);
int rank;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);

// parse command line arguments
if(argc<2){
    printf("usage: ./kaverage fileName.dat\n");
    return 0;
}

// initialize random seed
srand48(123456);

// read data from file
data_t data = dataReadWithClusterIds(argv[1]);

double start = mpiStartTimer();
// compute cluster centroids
clusters_t clusters = clusterSetupFromData(data);

// compute centroids
clusterComputeCentroids(clusters, data);
double time = mpiEndTimer(start);

if(rank == 0){
    printf("Size: %d Time: %f\n", size, time);
}

// output data
clusterOutput(clusters, "scatterdata.txt");

// finalize MPI
MPI_Finalize();

return 0;
}

```

1.2 clusters.c

```
#include <math.h>
#include <stdio.h>
#include <stdlib.h>

#include "data.h"
#include "clusters.h"

// set up random centroids for clusters
clusters_t clusterRandomSetup(int clusterK, data_t data){

    clusters_t clusters;

    clusters.K = clusterK;
    clusters.x = (float*) calloc(clusterK, sizeof(float));
    clusters.y = (float*) calloc(clusterK, sizeof(float));

    // work space for building new cluster centers
    clusters.tmpDegree = (int*) calloc(clusterK, sizeof(int));
    clusters.tmpx = (float*) calloc(clusterK, sizeof(float));
    clusters.tmpy = (float*) calloc(clusterK, sizeof(float));

    // initialize cluster centers
    for(int k=0;k<clusterK;++k){
        clusters.x[k] = drand48();
        clusters.y[k] = drand48();
    }

    // MPI buffer space for building new cluster centers
    clusters.bufDegree = (int*) calloc(clusterK, sizeof(int));
    clusters.bufx = (float*) calloc(clusterK, sizeof(float));
    clusters.buffy = (float*) calloc(clusterK, sizeof(float));

    return clusters;
}

// set up centroids from data
clusters_t clusterSetupFromData(data_t data){

    // find maximum cluster ID
    int clusterK = 0;
    for(int n=0;n<data.N;++n){
        int c = data.clusterIds[n];
        if(c>clusterK){
            clusterK = c;
        }
    }
}
```

```

    }
}
++clusterK;

// find maximum cluster idx
int maxK;
MPI_Allreduce(&clusterK, &maxK, 1, MPI_INT, MPI_MAX, MPI_COMM_WORLD);

clusters_t clusters;

clusters.K = maxK;
clusters.x = (float*) calloc(clusterK, sizeof(float));
clusters.y = (float*) calloc(clusterK, sizeof(float));

// work space for building new cluster centers
clusters.tmpDegree = (int*) calloc(clusterK, sizeof(int));
clusters.tmpx = (float*) calloc(clusterK, sizeof(float));
clusters.tmpy = (float*) calloc(clusterK, sizeof(float));

// MPI buffer space for building new cluster centers
clusters.bufDegree = (int*) calloc(clusterK, sizeof(int));
clusters.bufx = (float*) calloc(clusterK, sizeof(float));
clusters.buffy = (float*) calloc(clusterK, sizeof(float));

return clusters;
}

// compute centers of clusters
float clusterComputeCentroids(clusters_t clusters, data_t data){

    int K = clusters.K;

    float totalChange = 0;

    // zero accumulators for x,y,degree for all clusters
    for(int k=0;k<K;++k){
        clusters.bufx[k] = 0;
        clusters.buffy[k] = 0;
        clusters.bufDegree[k] = 0;
    }

    // sum up coordinates of all data points in each cluster
    for(int n=0;n<data.N;++n){
        int k = data.clusterIds[n];
        clusters.bufx[k] += data.x[n];

```

```

        clusters.bufo[k] += data.y[n];
        ++(clusters.bufoDegree[k]);
    }

    // sum up cluster x,y,counts globally
    MPI_Allreduce(clusters.bufo, clusters.tmpo, clusters.K, MPI_FLOAT, MPI_SUM, MPI_COMM_WORLD);
    MPI_Allreduce(clusters.bufo, clusters.tmpo, clusters.K, MPI_FLOAT, MPI_SUM, MPI_COMM_WORLD);
    MPI_Allreduce(clusters.bufoDegree, clusters.tmpDegree, clusters.K, MPI_FLOAT, MPI_SUM, MPI_COMM_WORLD);

    // compute centroids for clusters
    for(int k=0;k<clusters.K;++k){
        int degree = clusters.tmpDegree[k];
        if(degree>0){
            clusters.tmpo[k] /= (float) degree;
            clusters.tmpo[k] /= (float) degree;
        }else{
            // if empty cluster then choose random center
            clusters.tmpo[k] = drand48();
            clusters.tmpo[k] = drand48();
        }
    }

    // compute change in cluster centroid
    #if USE_POW==1
        float change =
            pow(clusters.x[k]-clusters.tmpo[k], 2) +
            pow(clusters.y[k]-clusters.tmpo[k], 2);
    #else
        float dx = clusters.x[k]-clusters.tmpo[k];
        float dy = clusters.y[k]-clusters.tmpo[k];

        float change = dx*dx+dy*dy;
    #endif

    totalChange += change;

    clusters.x[k] = clusters.tmpo[k];
    clusters.y[k] = clusters.tmpo[k];
}

return totalChange;
}

// label data points with cluster index
void clusterAssignDataPoints(clusters_t clusters, data_t data){

    //for each data point

```

```

for(int n=0;n<data.N;++n){

    // find its old cluster
    int oldk = data.clusterIds[n];

    float xn = data.x[n];
    float yn = data.y[n];
    data.clusterIds[n] = 0;

    #if USE_POW==1
        float minDistanceSquared = pow(xn-clusters.x[0],2) + pow(yn-clusters.y[0],2);
    #else
        float dx = xn-clusters.x[0];
        float dy = yn-clusters.y[0];
        float minDistanceSquared = dx*dx+dy*dy;
    #endif

    // find closest centroid
    for(int k=1;k<clusters.K;++k){

        #if USE_POW==1
            float distanceSquared = pow(xn-clusters.x[k],2) + pow(yn-clusters.y[k],2);
        #else
            dx = xn-clusters.x[k];
            dy = yn-clusters.y[k];
            float distanceSquared = dx*dx+dy*dy;
        #endif

        if(distanceSquared<minDistanceSquared){
            data.clusterIds[n] = k;
            minDistanceSquared = distanceSquared;
        }
    }
}

// output clusters to file from rank 0
void clusterOutput(clusters_t clusters, const char *fileName){

    int rank = 0;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if(rank==0){
        FILE *fp = fopen(fileName, "w");

        fprintf(fp, "x y cluster\n");
        for(int k=0;k<clusters.K;++k){
            fprintf(fp, "%f %f %d\n",

```

```

        clusters.x[k],
        clusters.y[k],
        k);
    }
    fclose(fp);
}
}

```

1.3 kmeans.c

```

#include <math.h>
#include <stdio.h>
#include <stdlib.h>

#include "omp.h"

#include "data.h"
#include "clusters.h"

// some comment

/*

    kmeans.c

    Purpose: read a data file and use kmeans clustering to assign data points to clusters

    Usage: ./kmeans dataFile.dat numberClusters

    To compile: gcc -o kmeans kmeans.c clusters.c data.c -lm

*/

// perform kmeans clustering using naive LLoyd's algorithm
int kmeans(data_t data, clusters_t clusters){

    float totalChange = 0;
    float tol = 1e-5;

    int iterations = 0;
    do{

        clusterAssignDataPoints(clusters, data);

        totalChange = clusterComputeCentroids(clusters, data);
    }
}

```

```

        ++iterations;

        printf("totalChange = %g\n", totalChange);

    }while(totalChange>tol*tol);

    printf("iterations = %d\n", iterations);

    return iterations;
}

// output data to file
void kmeansOutput(data_t data, clusters_t clusters, const char *fileName){

    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if(rank==0){
        FILE *fp = fopen(fileName, "w");

        fprintf(fp, "x y cluster\n");
        for(int k=0;k<data.N;++k){
            fprintf(fp, "%f %f %d\n",
                    data.x[k],
                    data.y[k],
                    data.clusterIds[k]+1);
        }
        fclose(fp);
    }
}

// perform kmeans clustering of data from file
int main(int argc, char **argv){

    // initialize MPI
    MPI_Init(&argc, &argv);

    // parse command line arguments
    if(argc<3){
        printf("usage: ./kmeans fileName.dat clusterCount\n");
        return 0;
    }

    int clusterK = atoi(argv[2]);

```



```

// initialize random seed
srand48(123456);

// read data from file
data_t data = dataRead(argv[1]);

// initialize cluster centroids
clusters_t clusters = clusterRandomSetup(clusterK, data);

// perform kmeans clustering
double tic = omp_get_wtime();
int iterations = kmeans(data, clusters);
double toc = omp_get_wtime();

printf("elapsed time per iteration = %3.2g\n", (toc-tic)/iterations);

// output data
kmeansOutput(data, clusters, "scatterdata.txt");

// shutdown MPI
MPI_Finalize();

return 0;
}

```

1.4 data.c

```

#include <math.h>
#include <stdio.h>
#include <stdlib.h>

#include "data.h"
#include "clusters.h"

// read data from file
data_t dataRead(const char *fileName){

    FILE *fp = fopen(fileName, "r");

    if(fp==NULL){
        printf("Failed to open file named: %s\n", fileName);
        exit(0);
    }

    int dataN;

```

```

data_t data;
char buf[BUFSIZ];

// read header
fgets(buf, BUFSIZ, fp);

// read number of data points
fgets(buf, BUFSIZ, fp);
sscanf(buf, "%d", &dataN);

// partition into chunks
int rank, size;
MPI_Comm_size(MPI_COMM_WORLD, &size);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
int chunkN = dataN/size;
int remainder = dataN - (chunkN*size);
data.N = chunkN;
if(rank<remainder) data.N += 1; // add 1 to low ranks to make sure have all data points

// allocate local space for data
data.x = (float*) calloc(data.N, sizeof(float));
data.y = (float*) calloc(data.N, sizeof(float));
data.clusterIds = (int*) calloc(data.N, sizeof(int));

// read section header
fgets(buf, BUFSIZ, fp);

// rank 0 reads all data from file and sends to other ranks
int tag = 999;
if(rank==0){

    // read own stuff
    for(int n=0;n<data.N;++n){
        fgets(buf, BUFSIZ, fp);
        sscanf(buf, "%f %f", data.x+n, data.y+n);
    }

    // buffer space for data to send
    float *bufx = (float*) calloc(chunkN+1, sizeof(float));
    float *bufy = (float*) calloc(chunkN+1, sizeof(float));

    // loop over ranks
    for(int r=1;r<size;++r){
        int sendN = (r<remainder) ? chunkN+1:chunkN;

        // read chunk to send to rank r

```

```

        for(int n=0;n<sendN;++n){
            fgets(buf, BUFSIZ, fp);
            sscanf(buf, "%f %f", bufx+n, bufy+n);
        }

        MPI_Send(bufx, sendN, MPI_FLOAT, r, tag+1, MPI_COMM_WORLD);
        MPI_Send(bufy, sendN, MPI_FLOAT, r, tag+2, MPI_COMM_WORLD);
    }

    // this rank is done with file
    fclose(fp);
}
else{
    // this rank is done with file
    fclose(fp);

    MPI_Status status;
    int recvN = data.N;
    int root = 0;
    MPI_Recv(data.x, recvN, MPI_FLOAT, root, tag+1, MPI_COMM_WORLD, &status);
    MPI_Recv(data.y, recvN, MPI_FLOAT, root, tag+2, MPI_COMM_WORLD, &status);
}

// print out a message from each rank
printf("Process rank %d owns %d data points\n", rank, data.N);

return data;
}

// read data from file
data_t dataReadWithClusterIds(const char *fileName){

    FILE *fp = fopen(fileName, "r");

    if(fp==NULL){
        printf("Failed to open file named: %s\n", fileName);
        exit(0);
    }

    int dataN;
    data_t data;
    char buf[BUFSIZ];

```

```

// read header
fgets(buf, BUFSIZ, fp);

// read number of data points
fgets(buf, BUFSIZ, fp);
sscanf(buf, "%d", &dataN);

// partition into chunks
int rank, size;
MPI_Comm_size(MPI_COMM_WORLD, &size);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
int chunkN = dataN/size;
int remainder = dataN - (chunkN*size);
data.N = chunkN;
if(rank<remainder) data.N += 1;

data.x = (float*) calloc(data.N, sizeof(float));
data.y = (float*) calloc(data.N, sizeof(float));
data.clusterIds = (int*) calloc(data.N, sizeof(int));

// read section header
fgets(buf, BUFSIZ, fp);

int tag = 999;
if(rank==0){

    // read own stuff
    for(int n=0;n<data.N;++n){
        fgets(buf, BUFSIZ, fp);
        sscanf(buf, "%f %f %d", data.x+n, data.y+n, data.clusterIds+n);
    }

    // buffer space for data to send
    float *bufx = (float*) calloc(chunkN+1, sizeof(float));
    float *bufy = (float*) calloc(chunkN+1, sizeof(float));
    int *bufClusterIds = (int*) calloc(chunkN+1, sizeof(int));

    // loop over ranks
    for(int r=1;r<size;++r){
        int sendN = (r<remainder) ? chunkN+1:chunkN;

        // read chunk to send to rank r
        for(int n=0;n<sendN;++n){
            fgets(buf, BUFSIZ, fp);
            sscanf(buf, "%f %f %d", bufx+n, bufy+n, bufClusterIds+n);

```

```

    }

    MPI_Send(bufx, sendN, MPI_FLOAT, r, tag+1, MPI_COMM_WORLD);
    MPI_Send(bufy, sendN, MPI_FLOAT, r, tag+2, MPI_COMM_WORLD);
    MPI_Send(bufClusterIds, sendN, MPI_INT, r, tag+3, MPI_COMM_WORLD);
}

// root done with file
fclose(fp);

}
else{
    // this rank done with file
    fclose(fp);

    MPI_Status status;
    int recvN = data.N;
    int root = 0;
    MPI_Recv(data.x, recvN, MPI_FLOAT, root, tag+1, MPI_COMM_WORLD, &status);
    MPI_Recv(data.y, recvN, MPI_FLOAT, root, tag+2, MPI_COMM_WORLD, &status);
    MPI_Recv(data.clusterIds, recvN, MPI_INT, root, tag+3, MPI_COMM_WORLD, &status);
}

// print out a message from each rank
printf("Process rank %d owns %d data points\n", rank, data.N);


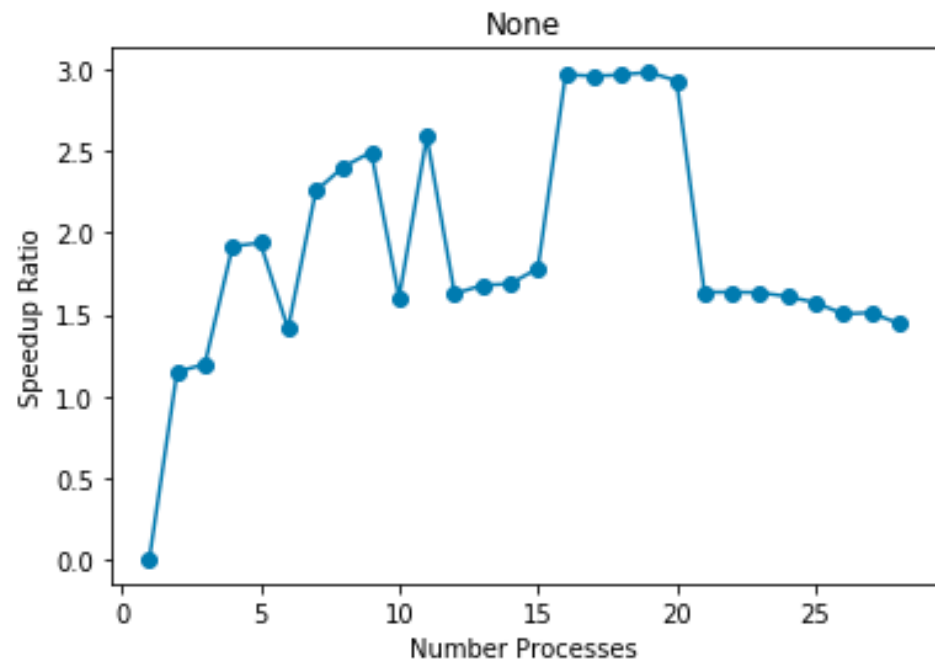
return data;
}



```

2 Task1

Optimization	Compiler	Runtime(sec)	Speedup
None	<u>mpicc</u>	0.009372	1
-O1	<u>mpicc</u>	0.002455	3.817515275
-O2	<u>mpicc</u>	0.002433	3.852034525
-O3	<u>mpicc</u>	0.002141	4.377393741
None	<u>icpc</u>	0.011382	1
-O1	<u>icpc</u>	0.002554	4.456538763
-O2	<u>icpc</u>	0.00219	5.197260274
-O3	<u>icpc</u>	0.002131	5.341154388

3 Task2





4 Kmeans chart

Optimization	Compiler	Runtime(sec)	Speedup
None	<u>mpicc</u>	0.72272	1
-O1	<u>mpicc</u>	0.37501	1.927201941
-O2	<u>mpicc</u>	0.356274	2.028551059
-O3	<u>mpicc</u>	0.407609	1.773071743
None	<u>icpc</u>	0.460384	1
-O1	<u>icpc</u>	0.186896	2.463316497
-O2	<u>icpc</u>	0.21147	2.177065305
-O3	<u>icpc</u>	0.210729	2.18472066

5 Discussion

I used exactly what was in the HW10 MPI example code to test the speedup ratios as well as the mpiStart and mpiStop time functions in the lesson 21 repositories. I couldn't figure out why the speedup ratio wouldn't reach higher than around 3. It seemed to speed up more or less as it was supposed to for 1-4 processes then plateau and actually start to get slower as more processes were added.